

CS80S - Theory and practice of secure systems

Fall 2024

Tue / Thu : 2 pm to 3:30 pm

Shravan Narayan

UT Austin

Announcements

Sign up to present a paper

Minor changes to papers. Please check what you've signed up for!

Think about your projects! Did you figure out compute resources?

Writeups submission links will be posted later today

Writeups and attendance starting Tuesday

Today

1. How to present a paper
2. Background on memory safety
3. Background on side-channels
4. Security News

Paper presentation

A good resource

[\(An Opinionated Talk\) On Preparing Good Talks](#) – Ranjit Jhala (2018)

Paper presentation: preparation

1. Watch the talk from the previous slide
2. Read the paper
3. Repeat step 1 until comfortable
4. Skim papers that are cited by the paper
5. Skim papers that have cited your paper (Use Google Scholar)
6. Watch any videos from authors/others presenting this work
7. Concisely state the high-level problem and the solution.
8. Shortlist important points – Do not try to cover everything :)

Come talk to me if you need help understanding the paper.

Paper presentation: Slides

1. You can make slides or give a whiteboard talk
2. Plan for a 30 min talk: typical average is 1 to 2 slides per minute
3. Make sure the talk clarifies
 - What the problem is
 - Why the problem matters
 - Why existing solutions won't help
 - What is the new approach proposed by authors
 - Relevant implementation details that can fit within the time of the talk
 - Evaluation of the new idea
4. You can borrow slides/figures from the paper if you cite sources
 - Original slides for paper are usually too dense

Background on memory safety

Why does memory safety matter?

From Pearl to Pegasus Bahraini Government Hacks Activists with NSO Group Zero-Click iPhone Exploits

By Bill Marczak, Ali Abdulemam¹, Noura Al-Jizawi, Siena Anstis, Kristin Berdan, John Scott-Railton, and Ron Deibert

[1] Red Line for Gulf

August 24, 2021

Summary & Key Findings

- We identified nine Bahraini activists whose iPhones were successfully hacked with NSO Group's Pegasus spyware between June 2020 and February 2021. Some of the activists were hacked using two zero-click iMessage exploits: [the 2020 KISMET exploit](#) and a 2021 exploit that we call **FORCEDENTRY**.

About the Citizen Lab

The Citizen Lab is an interdisciplinary laboratory based at the Munk School of Global Affairs & Public Policy, University of Toronto, focusing on research, development, and high-level strategic policy and legal engagement at the intersection of information and communication technologies, human rights, and global security.

We use a “mixed methods” approach to research combining practices from political science, law, computer science, and area studies. Our research includes: investigating digital es-

<https://citizenlab.ca/2021/08/bahrain-hacks-activists-with-nso-group-zero-click-iphone-exploits/>

Why does memory safety matter?

When the **FORCEDENTRY** exploit was being fired at a device, the device logs showed crashes associated with *IMTranscoderAgent*. The crashes appeared to be segfaults generated by invoking the *copyGifFromPath:toDestinationPath:error* function on files received via iMessage.

The crashes appeared to be of two types. Type one crashes indicate that the chain of events set off by invoking *copyGifFromPath:toDestinationPath:error* ultimately crashed while apparently invoking ImageIO's functionality for rendering Adobe Photoshop PSD data.

```
Thread 1 name: Dispatch queue: com.apple.root.default-qos
Thread 1 Crashed:
0: ImageIO 0x181b326f0 __ZN13PSDReadPlugin12GetRangeInfoEP9LayerInfoP19IIIOImageReadSessionRK13PSDPluginData + 244
1: ImageIO 0x181b326ec __ZN13PSDReadPlugin12GetRangeInfoEP9LayerInfoP19IIIOImageReadSessionRK13PSDPluginData + 240
2: ImageIO 0x181b32abc __ZN13PSDReadPlugin11DecodeLayerEP9LayerInfoRK6CGRectP19IIIOImageReadSessionRK14ReadPluginDataRK13PSDPluginData + 448
3: ImageIO 0x181b32d40 __ZN13PSDReadPlugin11MergeLayersEP19IIIOImageReadSessionRK20IIIODecodeFrameParamsRK14ReadPluginDataRK13PSDPluginData + 276
4: ImageIO 0x181b33248 __ZN13PSDReadPlugin11DecodeBlockEP19IIIOImageReadSessionRK20IIIODecodeFrameParamsRK14ReadPluginDataRK13PSDPluginData + 232
5: ImageIO 0x181b34c4c __ZN13PSDReadPlugin12DecodeBlocksEP12IIIOImageReadRK14ReadPluginDataRK13PSDPluginDataRNS3__16vectorI20IIIODecodeFrameParamsNS8_9allocatorISA_EEEE_block_invoke + 116
6: libdispatch.dylib 0x18004b860 __dispatch_client_callout2 + 20
7: libdispatch.dylib 0x18005f988 __dispatch_apply_serial + 120
8: libdispatch.dylib 0x18004b81c __dispatch_client_callout + 20
9: libdispatch.dylib 0x180050c7c __dispatch_sync_function_invoke + 56
10: libdispatch.dylib 0x18005f944 __dispatch_apply_f + 908
11: ImageIO 0x181b34bcc __ZN13PSDReadPlugin120codeBlocksEP12IIIOImageReadRK14ReadPluginDataRK13PSDPluginDataRNS3__16vectorI20IIIODecodeFrameParamsNS8_9allocatorISA_EEEE + 140
12: ImageIO 0x181b34364 __ZN13PSDReadPlugin17copyImageBlockSetEP7InfoRecP15CGImageProvider6CGRect6CGSizePK14_CFDictionary + 1028
13: ImageIO 0x181a28388 __ZN10IIIO_Reader21CopyImageBlockSetProcEPVP15CGImageProvider6CGRect6CGSizePK14_CFDictionary + 152
14: ImageIO 0x181a36510 __ZN20IIIOImageProviderInfo28copyImageBlockSetWithOptionsEP15CGImageProvider6CGRect6CGSizePK14_CFDictionary + 756
15: ImageIO 0x181a33e74 __ZN20IIIOImageProviderInfo28CopyImageBlockSetWithOptionsEPVP15CGImageProvider6CGRect6CGSizePK14_CFDictionary + 584
16: CoreGraphics 0x182004c68 imageProvider_retain_data + 92
17: CoreGraphics 0x181decac8 CGDataProviderRetainData + 88
18: CoreGraphics 0x181ff1e04 CGAccessSessionCreate + 108
19: CoreGraphics 0x181e21bc8 CGDataProviderCopyData + 168
20: CoreGraphics 0x181d95230 CGImageGetDataProviderInternal + 268
21: CoreGraphics 0x181eccc10 __img_image + 600
22: CoreGraphics 0x181ecc6e8 CGSImageDataLock + 1004
23: CoreGraphics 0x181ceb2bc __pipe_AcquireRIPImageData + 716
24: CoreGraphics 0x181ee25c __pipe_DrawImage + 1152
25: CoreGraphics 0x181ed3b40 CGContextDrawImageWithOptions + 1216
26: ImageIO 0x181ae2730 __ZN14GIFWritePlugin10writeSingleFrameEv + 1016
27: ImageIO 0x181ae3890 __ZN14GIFWritePlugin8writeAllEv + 500
28: ImageIO 0x181ad76c __ZN14IIIO_Writer_GIF8writeEPvS0 + 36
29: ImageIO 0x181ae654 __ZN19IIIOImageDestination19finalizeDestinationEv + 548
30: ImageIO 0x181aeffd4 CGImageDestinationFinalize + 132
31: IMSharedUtilities 0x18f77c9b4 writeNewFileAtPath:withProperties:fromImageSource:error: + 236
32: IMSharedUtilities 0x18f77cd88 copyGifFromPath:toDestinationPath:error: + 348
33: IMTranscoderAgent 0xf08e58c8
```

<https://citizenlab.ca/2021/08/bahrain-hacks-activists-with-nso-group-zero-click-iphone-exploits/>

Goal – Before next class, you should know

1. What is a memory safety attack?
2. How does it compromise the system?

But first: basic OS concepts

Model of process – stack, heap, code, functions calls

Virtual memory vs physical memory

Memory pages, page permissions, and segmentation faults

But first: basic OS concepts

Model of process – stack, heap, code, functions calls

Virtual memory vs physical memory

Memory pages, page permissions, and segmentation faults

C is a language to target these abstractions, but...

C/C++ is not “memory safe”

Designed > 30 years ago when performance was key

⇒ Not secure by default

A problem for old and new software alike

Memory safety root cause: Missing checks in code

- Buffer overflows
- Pointer dereference out-of-bounds
- Use-after-free
- Type confusion

Google Chrome: ~70% of bugs (2015–2020)




Image from the Chromium project blog

<https://www.chromium.org/Home/chromium-security/memory-safety/>

C/C++ is not “memory safe”

While a programmer must fix all possible bugs and defenses must cover all possible attack vectors to secure a system, an adversary only needs to find one exploitable bug to gain control of the system.

Memory safety

Classically divided into a few sub-areas

- Spatial safety
- Temporal safety
- Type safety

Memory safety: spatial safety violation

```
int main(int argc, char** argv) {  
    char buffer[100];  
  
    for(int i = 0; i <= 100; i++) {  
        buffer[i] = 'a';  
        // buffer[100] is OOB, buffer overflow, memory unsafe, spatially unsafe  
    }  
  
    // Extra question: this is C UB, how does this related to memory safety  
}
```


Memory safety: temporal safety violation

```
int main(int argc, char** argv) {  
    char* buffer = malloc(16);  
  
    char* copy = buffer;  
    free(copy);  
  
    // violation as buffer no longer points to a valid object  
    buffer[12] = 23;  
}
```

Memory safety: type safety violation

```
class B {  
public:  
    int b_field;  
};  
  
class D {  
public:  
    virtual void d_func() {}  
    int d_field;  
};
```

```
int main(int argc, char** argv) {  
    D* d_ptr = new D;  
  
    // Bad cast  
    B* b_ptr = (B*) d_ptr;  
  
    // Type confusion!  
    b_ptr->b_field = 0x43;  
    b_ptr->d_func();  
}
```

Mem safety attacks break program properties

1. Memory integrity

- Data integrity
- Data pointer integrity
- Code pointer integrity

2. Memory confidentiality

- Leak data
- Leak location of code

Mem safety attacks break program properties

1. Memory integrity

- Data integrity
- Data pointer integrity
- Code pointer integrity

2. Memory confidentiality

- Leak data
- Leak location of code

30.2	A Selection of Low-Level Attacks on C Software	635
30.2.1	Attack 1: Corruption of a Function Return Address on the Stack	636
30.2.2	Attack 2: Corruption of Function Pointers Stored in the Heap	638
30.2.3	Attack 3: Execution of Existing Code via Corrupt Pointers	640
30.2.4	Attack 4: Corruption of Data Values that Determine Behavior	643

Background on side-channels

Transient-Execution Attacks: A Computer Architect Perspective

Luís Fiolhais, Leonel Sousa (2023)

Normal attacks

Learn data due to bugs in hw/sw (buffer overflows)

Side-channel attacks

Learn data by measuring indirect effects of hw/sw impls. (caches)

Architectural concepts

Translation lookaside buffer (TLB) and page table entries (PTE)

Superscalar, Instruction-level parallelism (ILP)

Out of order execution (OOO), Reorder buffer (ROB)

Branch predictors & Speculative execution

Instruction timing and microcode

Store to load forwarding

Frequency scaling

Caches and the cache hierarchy

- Ways and sets
- Inclusive vs Exclusive
- Writeback vs. Write-through

Prime + Probe cache attack

```
// I am sharing a single core computer with a webserver
```

```
void on_http_request(int user_id) { // Server code
```

```
    auto userinfo = user_data[user_id];
```

```
    // Some IO
```

```
}
```

```
void prime_and_probe() { // My code
```

```
    // 1. Create an array the size of the cache
```

```
    // 2. Access the entire array to fill up the cache and then sleep
```

```
    // 3. On wakeup, measure time to access each array element
```

```
    // 4. Slow-elements correspond to cache collisions
```

```
}
```

Question 1:

How can I use this to learn information

Question 2:

What are the ways this is not realistic?

Security News

Let's look at some news

Windows Patch Tuesdays

<https://www.bleepingcomputer.com/news/microsoft/microsoft-august-2024-patch-tuesday-fixes-9-zero-days-6-exploited/>

Google Chrome bugs

<https://thehackernews.com/2024/08/google-fixes-high-severity-chrome-flaw.html>